

▼ Heart Attack Prediction

The main aim of this project is to predict the chances of heart attack by analysing some medical properties like blood pressure,chest pain,sugar,cholesterol,maximum heart rate achieved

```
#Importing all the necessary librarys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense,Dropout
from sklearn.metrics import classification_report,accuracy_score,roc_curve,roc_auc_score

#Loading dataset and Reading
df = pd.read_csv('heart.csv')
df.head()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	tha
0	63	1	3	145	233	1	0	150	0	2.3	0	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    age         303 non-null    int64
1    sex         303 non-null    int64
2    cp          303 non-null    int64
3    trtbps      303 non-null    int64
4    chol        303 non-null    int64
5    fbs         303 non-null    int64
6    restecg     303 non-null    int64
7    thalachh    303 non-null    int64
8    exng        303 non-null    int64
9    oldpeak     303 non-null    float64
10   slp         303 non-null    int64
11   caa         303 non-null    int64
12   thall       303 non-null    int64
13   output      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

▼ Data Preprocessing

```
df.nunique()
```

```
age      41
sex       2
cp        4
trtbps   49
```

```
chol      152
fbs       2
restecg   3
thalachh  91
exng      2
oldpeak   40
slp       3
caa       5
thall     4
output    2
dtype: int64
```

We have 5 numerical feature and 8 categorical feature in our data set

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trtbps      303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalachh    303 non-null    int64
8   exng        303 non-null    int64
9   oldpeak     303 non-null    float64
10  slp         303 non-null    int64
11  caa         303 non-null    int64
12  thall       303 non-null    int64
13  output      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

All our categorical nature column has integer type of data type so first we convert into an object data type

```
cat_columns = ['sex', 'cp', 'fbs', 'restecg', 'exng', 'slp', 'caa', 'thall', 'output']
num_columns = ['age', 'trtbps', 'oldpeak', 'chol', 'thalachh']
df[cat_columns] = df[cat_columns].astype(str)
```

```
df.describe()
```

	age	trtbps	chol	thalachh	oldpeak
count	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	131.623762	246.264026	149.646865	1.039604
std	9.082101	17.538143	51.830751	22.905161	1.161075
min	29.000000	94.000000	126.000000	71.000000	0.000000
25%	47.500000	120.000000	211.000000	133.500000	0.000000
50%	55.000000	130.000000	240.000000	153.000000	0.800000
75%	61.000000	140.000000	274.500000	166.000000	1.600000
max	77.000000	200.000000	564.000000	202.000000	6.200000

Median and mean values of all the numerical features are comparable. Therefore, I don't think there are any outliers in our numerical features

▼ Handling Missing values

```
df.isnull().sum()
```

```

age      0
sex      0
cp       0
trtbps   0
chol     0
fbs      0
restecg  0
thalachh 0
exng     0
oldpeak  0
slp      0
caa      0
thall    0
output   0
dtype: int64

```

As there is no null value our data is clean for preprocessing

Visualization

Let us check how our numerical features change based on target columns

```

sns.set_context('notebook', font_scale=1.2)
fig, ax=plt.subplots(2,2,figsize=(20,20))
plt.title('Distribution of various feature based on target ')
a1 = sns.histplot(x='age', data= df, hue= 'output',kde=True, ax=ax[0, 0], palette='winter')
a1.set(xlabel = 'Age', title= 'Distribution of age based on target variable')

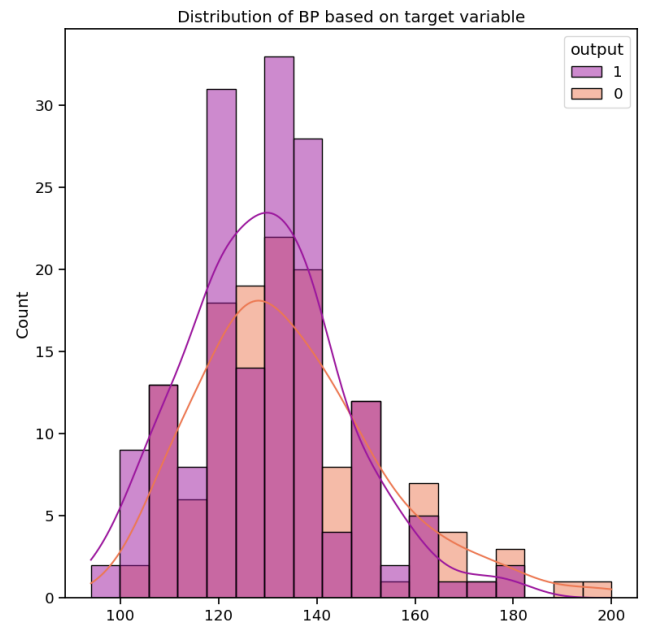
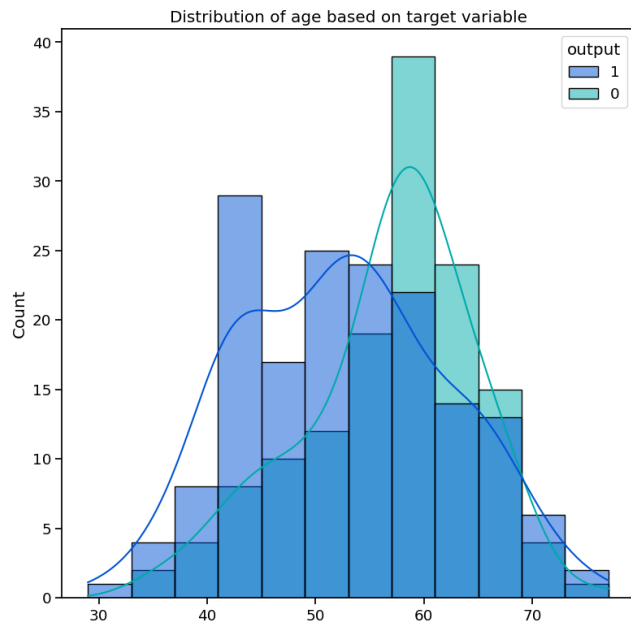
a2 = sns.histplot(x='trtbps', data= df, hue= 'output', kde= True, ax= ax[0, 1], palette='plasma')
a2.set(xlabel = 'Resting blood pressure (in mm Hg)', title= 'Distribution of BP based on target variable')

a3 = sns.histplot(x='chol', data= df, hue= 'output', kde= True, ax= ax[1, 0], palette='winter')
a3.set(xlabel = 'Cholestoral in mg/dl', title= 'Distribution of Cholestrol based on target variable')

a4 = sns.histplot(x='thalachh', data=df, hue= 'output', kde= True, ax= ax[1, 1], palette='plasma')
a4.set(xlabel = 'Max Heart Rate Achieved', title= 'Distribution of maximum heart rate achieved based on target variable')

plt.show()

```



We can clearly seen that the person who had maximum heart rate achieved having a great chance of heart attack

Distribution of Cholestrol based on target variable

Distribution of maximum heart rate achieved based on target variable

To check outlier is present or not

```
sns.set_context('notebook', font_scale= 1.2)
fig, ax = plt.subplots(2, 2, figsize = (20, 10))

plt.suptitle('Boxplot of various features based on target variable', fontsize = 20)

b1 = sns.boxplot(x = 'age', data=df, ax= ax[0, 0], color = '#40bf80')
b1.set(xlabel = 'Age')

b2 = sns.boxplot(x = 'trtbps', data=df, ax= ax[0, 1], color='#40bf80')
b2.set(xlabel = 'Resting blood pressure (in mm Hg)')

b3 = sns.boxplot(x = 'chol', data=df, ax= ax[1, 0], color= '#40bf80')
b3.set(xlabel = 'Cholestoral in mg/dl')

b4 = sns.boxplot(x = 'thalachh', data=df, ax= ax[1, 1], color = '#40bf80')
b4.set(xlabel = 'Max Heart Rate Achieved')

plt.show()
```

Boxplot of various features based on target variable




As we clearly seen that there are outlier present in cholesterol and blood pressure column.



```
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True)
```

<Axes: >




As with the help of heatmap we see that there is no correlation in our numerical columns

```
df['output'].value_counts()
```

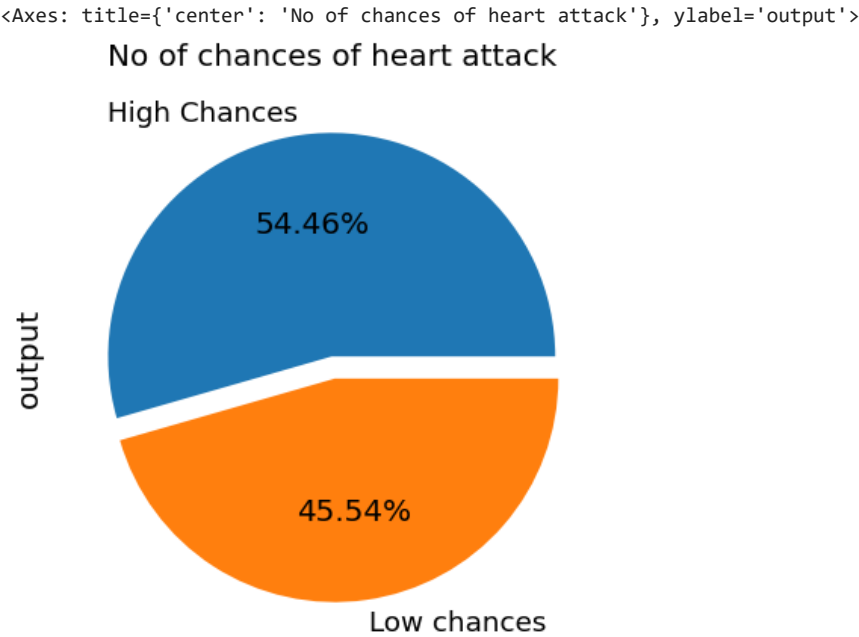
```
1    165
0    138
Name: output, dtype: int64
```

```
sns.countplot(data=df,x=df['output'])
```



Hence with above countplot we can clearly seen that our target column is balanced in nature.

```
plt.title('No of chances of heart attack')
df['output'].value_counts().plot.pie(autopct='%1.2f%',labels=['High Chances','Low chances'],explode=(0,0.1))
```



```
df.head()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	

```
cat_columns = ['sex', 'cp', 'fbs', 'restecg', 'exng', 'slp', 'caa', 'thall', 'output']
num_columns = ['age', 'trtbps', 'oldpeak', 'chol', 'thalachh']
df[cat_columns] = df[cat_columns].astype(float)
```

```
#To seperate x and y
x=df.iloc[:, :-1].values
```

```
array([[57., 0., 1., ..., 1., 2., 0.],
       [57., 1., 0., ..., 1., 3., 0.],
       [68., 1., 0., ..., 2., 3., 0.],
       ...,
       [41., 0., 1., ..., 0., 2., 1.],
       [37., 1., 2., ..., 0., 2., 1.],
       [63., 1., 3., ..., 0., 1., 1.]])
```

[illegible]

```
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.20,random_state=1)
```

▼ Early Stopping

```
#Early stopping
from tensorflow.keras.callbacks import EarlyStopping
early_stop=EarlyStopping(monitor="val_loss",mode="min",verbose=1,patience=10)
```

```
ann=Sequential()  
#Hidden Layer  
ann.add(Dense(units=100,activation="relu"))  
ann.add(Dropout(0.20))  
ann.add(Dense(units=50,activation="relu"))  
#ann.add(Dense(units=100,activation="relu"))  
#ann.add(Dropout(0.20))  
#ann.add(Dense(units=50,activation="relu"))  
  
ann.add(Dense(units=1,activation="sigmoid"))
```

```
ann.compile(optimizer="adam",loss="binary_crossentropy",metrics=["accuracy"])
```


```
ann.fit(xtrain,ytrain,validation_data=(xtest,ytest),verbose=1,callbacks=[early_stop],batch_size=40,epochs=100)
```

```
=====] - 0s 9ms/step - loss: 0.3444 - accuracy: 0.9091 - val_loss: 0.3712 - val_accuracy: 0.9091
8/100
=====] - 0s 9ms/step - loss: 0.3239 - accuracy: 0.9008 - val_loss: 0.3525 - val_accuracy: 0.9008
9/100
=====] - 0s 9ms/step - loss: 0.3074 - accuracy: 0.9091 - val_loss: 0.3399 - val_accuracy: 0.9091
10/100
=====] - 0s 8ms/step - loss: 0.2833 - accuracy: 0.9215 - val_loss: 0.3307 - val_accuracy: 0.9215
11/100
=====] - 0s 13ms/step - loss: 0.2873 - accuracy: 0.9132 - val_loss: 0.3258 - val_accuracy: 0.9132
12/100
=====] - 0s 12ms/step - loss: 0.2596 - accuracy: 0.9256 - val_loss: 0.3193 - val_accuracy: 0.9256
13/100
=====] - 0s 11ms/step - loss: 0.2459 - accuracy: 0.9298 - val_loss: 0.3117 - val_accuracy: 0.9298
14/100
=====] - 0s 9ms/step - loss: 0.2392 - accuracy: 0.9256 - val_loss: 0.3059 - val_accuracy: 0.9256
15/100
=====] - 0s 15ms/step - loss: 0.2457 - accuracy: 0.9174 - val_loss: 0.3064 - val_accuracy: 0.9174
16/100
=====] - 0s 10ms/step - loss: 0.2196 - accuracy: 0.9256 - val_loss: 0.3170 - val_accuracy: 0.9256
17/100
=====] - 0s 10ms/step - loss: 0.2205 - accuracy: 0.9256 - val_loss: 0.3246 - val_accuracy: 0.9256
18/100
=====] - 0s 11ms/step - loss: 0.2259 - accuracy: 0.9256 - val_loss: 0.3229 - val_accuracy: 0.9256
19/100
=====] - 0s 9ms/step - loss: 0.2137 - accuracy: 0.9215 - val_loss: 0.3179 - val_accuracy: 0.9215
20/100
=====] - 0s 9ms/step - loss: 0.2040 - accuracy: 0.9256 - val_loss: 0.3102 - val_accuracy: 0.9256
21/100
=====] - 0s 9ms/step - loss: 0.1986 - accuracy: 0.9339 - val_loss: 0.2973 - val_accuracy: 0.9339
22/100
=====] - 0s 13ms/step - loss: 0.1961 - accuracy: 0.9298 - val_loss: 0.2941 - val_accuracy: 0.9298
23/100
=====] - 0s 9ms/step - loss: 0.1846 - accuracy: 0.9256 - val_loss: 0.2894 - val_accuracy: 0.9256
24/100
=====] - 0s 11ms/step - loss: 0.1816 - accuracy: 0.9421 - val_loss: 0.2766 - val_accuracy: 0.9421
25/100
=====] - 0s 8ms/step - loss: 0.1808 - accuracy: 0.9256 - val_loss: 0.2808 - val_accuracy: 0.9256
26/100
=====] - 0s 14ms/step - loss: 0.1988 - accuracy: 0.9298 - val_loss: 0.3036 - val_accuracy: 0.9298
27/100
=====] - 0s 10ms/step - loss: 0.1779 - accuracy: 0.9215 - val_loss: 0.3107 - val_accuracy: 0.9215
28/100
=====] - 0s 9ms/step - loss: 0.1697 - accuracy: 0.9298 - val_loss: 0.3031 - val_accuracy: 0.9298
29/100
=====] - 0s 11ms/step - loss: 0.1757 - accuracy: 0.9339 - val_loss: 0.2844 - val_accuracy: 0.9339
30/100
=====] - 0s 8ms/step - loss: 0.1715 - accuracy: 0.9339 - val_loss: 0.2806 - val_accuracy: 0.9339
31/100
=====] - 0s 9ms/step - loss: 0.1627 - accuracy: 0.9298 - val_loss: 0.2797 - val_accuracy: 0.9298
32/100
=====] - 0s 11ms/step - loss: 0.1680 - accuracy: 0.9463 - val_loss: 0.2889 - val_accuracy: 0.9463
33/100
=====] - 0s 11ms/step - loss: 0.1663 - accuracy: 0.9504 - val_loss: 0.2960 - val_accuracy: 0.9504
34/100
=====] - 0s 9ms/step - loss: 0.1654 - accuracy: 0.9380 - val_loss: 0.3005 - val_accuracy: 0.9380
34: early stopping
.callbacks.History at 0x7f30b2bf9b50>
```

```
ann.history.history #data points of loss with respect to epochs
```


[illegible]

```
plt.plot(pd.DataFrame(ann.history.history))
plt.show()
```



```
ypred=ann.predict(xtest)
```

```
2/2 [=====] - 0s 8ms/step
```

```
array([[0],
```

	precision	recall	f1-score	support
0.0	0.83	0.97	0.89	30
1.0	0.96	0.81	0.88	31
accuracy			0.89	61
macro avg	0.90	0.89	0.88	61
weighted avg	0.90	0.89	0.88	61

https://colab.research.google.com/drive/1JLoDNW3j4YLUtXoGefHPm7iRfmC-1aZ5#scrollTo=DDDq1MXME_ck&printMode=true

```
[[29  1]  
 [ 6 25]]
```

As you can clearly seen by our matrix that our model classification is good


▼ Roc Auc Curve

```
from sklearn.metrics import roc_curve,roc_auc_score
```

```
fpr, tpr, thresholds = roc_curve(ytest,ypred)
```

```
plt.figure(figsize=(6,4))  
plt.plot(fpr, tpr, linewidth=2, color= 'teal')  
plt.plot([0,1], [0,1], 'r--' )  
plt.title('ROC Curve')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')
```

```
plt.show()
```



```
print(roc_auc_score(ytest,ypred))
```

```
0.8865591397849462
```

Here you can see that our roc auc curve is closest to 1,Hence our model is performing good while doing prediction whether the person has heart attack chances or not.

✓ 0s completed at 12:33

● ✕